

<!-- Awin verification 001 -->

[Skip to content](#)

**Novixys Software Dev Blog**

MENU

# How to Setup an Apache Hadoop Cluster on AWS EC2

This article demonstrates how to get Apache Hadoop up and running on an Amazon EC2 cluster.



## Contents [\[hide\]](#)

- [Introduction](#)
- [Pre-Requisites](#)
- [AWS EC2 Startup](#)
  - [Select Instance](#)
  - [Instance Type](#)
  - [Instance Details](#)
  - [Storage](#)
  - [Instance Tags](#)
  - [Security Group](#)
  - [Launch Instances](#)

- [Naming the Instances](#)
- [Setting Up Instances](#)
  - [Copy Instance Public DNS Name](#)
- [Common Setup on All Nodes](#)
  - [All Nodes: Install Java](#)
  - [All Nodes: Install Apache Hadoop](#)
  - [All Nodes: Setup JAVA\\_HOME](#)
  - [All Nodes: Update core\\_site.xml](#)
  - [All Nodes: Create Data Dir](#)
- [Configuring NameNode](#)
  - [Namenode: Password Less SSH](#)
  - [Datanodes: Setup Public Key](#)
  - [Namenode: Setup SSH Config](#)
  - [Namenode: Setup HDFS Properties](#)
  - [Namenode: Setup MapReduce Properties](#)
  - [Namenode: Setup YARN Properties](#)
  - [Namenode: Setup Master and Slaves](#)
- [Configuring Data Nodes](#)
- [Starting the Hadoop Cluster](#)
  - [Check the Web UI](#)
- [Summary](#)
  - [See Also](#)

## Introduction

Lets talk about how to setup an [Apache Hadoop](#) cluster on [AWS](#).

In a [previous article](#), we discussed setting up a Hadoop processing pipeline on a single node (laptop). That involved running all the components of Hadoop on a single machine. In the setup we discuss here, we setup a multi-node cluster to run processing jobs.

Our setup involves a single NameNode and three DataNodes which serve as processing slaves.

Starting with setting up the AWS EC2 resources, we take you all the way through complete configuration of the machines in this arrangement.

We use Apache Hadoop 2.7.3 for this demonstration.

## Pre-Requisites

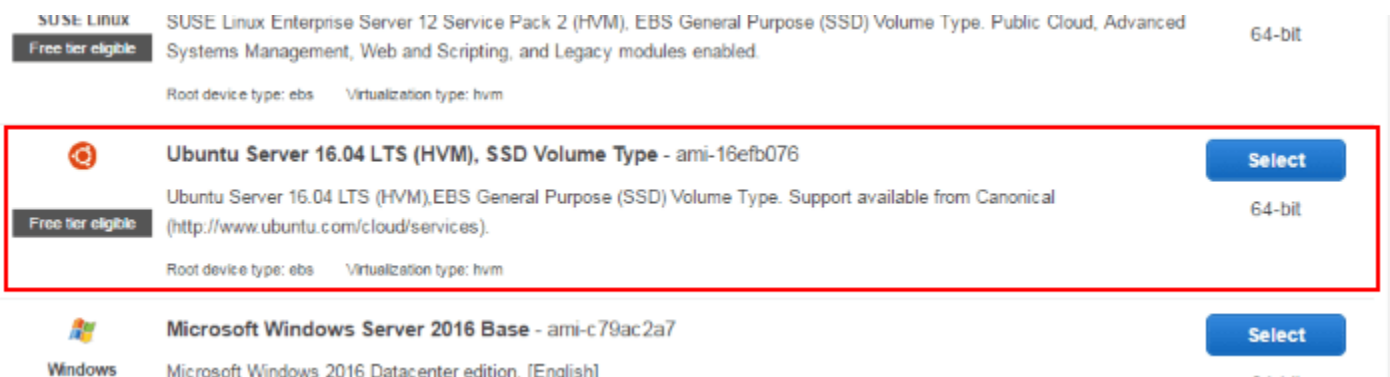
Sign up for an AWS account if you don't already have one. You get some resources free for the first year, including an EC2 Micro Instance.

## AWS EC2 Startup

We will now create 4 instances of Ubuntu Server 16.04 LTS using Amazon EC2.

### Select Instance

Go to your AWS Console, Click on Launch Instance and select Ubuntu Server 16.04 LTS.



### Instance Type

For the instance type, we choose **t2.micro** since that is sufficient for the purposes of the demo. If you have a need for a high-memory or high-cpu instance, you can select one of those.

<input type="checkbox"/>	General purpose	t2.nano	1	0.5	EBS on
<input checked="" type="checkbox"/>	General purpose	t2.micro Free tier eligible	1	1	EBS on
<input type="checkbox"/>	General purpose	t2.small	1	2	EBS on
<input type="checkbox"/>	General purpose	t2.medium	2	4	EBS on

Click Next to Configure Instance Details

## **Instance Details**

Here, we request 4 instances of the selected machine type. We also choose a subnet (us-west-1b) just so we can launch into the same location if we need more machines.

Click Next to Add Storage

## **Storage**

For our purpose, the default instance storage of 8GB is sufficient. If you need more storage, either increase the size or attach a disk by clicking “Add Volume”. If you add a volume, you will need to attach the volume to your instance, format it and mount it. Since this is a beginner tutorial, these steps are not covered here.

Volume Type ⓘ	Device ⓘ	Snapshot ⓘ	Size (GiB) ⓘ	Volume Type ⓘ	IOPS ⓘ	Throughput (MB/s) ⓘ
Root	/dev/sda1	snap- 0d553b777dbb9ce13	<input type="text" value="8"/>	<input type="text" value="General Purpose SSD (GP2)"/>	100 / 3000	N/A

[Add New Volume](#)

Click Next to Add Tags to your instances.

## Instance Tags

A tag allows you to identify your instance with a name you can choose.

Click Add Tag, set the Key to “Name” and value to “Hadoop”. We will use this tag to re-label our instances as “namenode”, “datanode1” and so on later on. For now leave the value of all the instances as “Hadoop”.

Key (127 characters maximum)	Value (255 characters maximum)
<input type="text" value="Name"/>	<input type="text" value="Hadoop"/>

[Add another tag](#) (Up to 50 tags maximum)

Click Next to configure Security Group for the instances.

## Security Group


For the security group, we create a completely open security group for the purposes of testing.

Assign a security group: ☒ Create a new security group  
☐ Select an existing security group

Security group name:

Description:

Type ⓘ	Protocol ⓘ	Port Range ⓘ	Source ⓘ
All traffic ▼	All	0 - 65535	Anywhere ▼ 0.0.0.0/0, ::/0

 **Warning**  
Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Finally we get to the Launch screen.

## Launch Instances

Review the information again and click Launch to start the instances. You will need to create a key pair or use an existing key pair. Follow the instructions on the Launch Wizard to create a new key pair. Then click “I acknowledge ..” and then Launch Instances.

## Select an existing key pair or create a new key pair



A key pair consists of a **public key** that AWS stores, and a **private key file** that you store. Together, they allow you to connect to your instance securely. For Windows AMIs, the private key file is required to obtain the password used to log into your instance. For Linux AMIs, the private key file allows you to securely SSH into your instance.

Note: The selected key pair will be added to the set of keys authorized for this instance. Learn more about [removing existing key pairs from a public AMI](#).

Choose an existing key pair ▼

### Select a key pair

Hadoop ▼

☒ I acknowledge that I have access to the selected private key file (Hadoop.pem), and that without this file, I won't be able to log into my instance.

Cancel


Launch Instances

You can now go to the instances page and check on the status of the instances.

## Naming the Instances

On the instances page, let us setup the names of the instances. These are not DNS names, but names we assign to help us distinguish between them.

Click the pencil icon next to the name and setup the names as shown.

<div> <div>Launch Instance</div> <div>Connect</div> <div>Actions ▾</div> </div>				
<div> <div> <div></div> <div>Filter by tags and attributes or search by keyword</div> </div> </div>				
<input type="checkbox"/>	Name ▾	Instance ID ▴	Instance Type ▾	Availability Zone ▾
<input checked="" type="checkbox"/>	Hadoop-datanode3	i-02bfd96c809300db0	t2.micro	us-west-1b
<input type="checkbox"/>	Hadoop-datanode2	i-068b63a7c8770b6b4	t2.micro	us-west-1b
<input type="checkbox"/>	Hadoop-datanode1 	i-0b02575e7c564c3d7	t2.micro	us-west-1b
<input type="checkbox"/>	Hadoop-namenode	i-0eda579a312589c2e	t2.micro	us-west-1b

## Setting Up Instances

Once the instances are up and running, it is time to set them up for our purpose. This includes the following:

- Setup password-less login between the namenode and the datanodes.
- Install java.
- Setup Hadoop.

### Copy Instance Public DNS Name

We now need to copy the Public DNS Name of each node (1 namenode and 3 datanodes). These names are used in the configuration steps below. Since the DNS is specific to each setup, we refer to the names as follows.

For example, in the description below, if you see <nnode>, substitute with the value of <NameNode Public DNS>. Similarly for <dnode1> and so on.

nnode	<NameNode Public DNS>
dnode1	<DataNode1 Public DNS>



dnnode2	<DataNode2 Public DNS>
dnnode3	<DataNode3 Public DNS>

## Common Setup on All Nodes

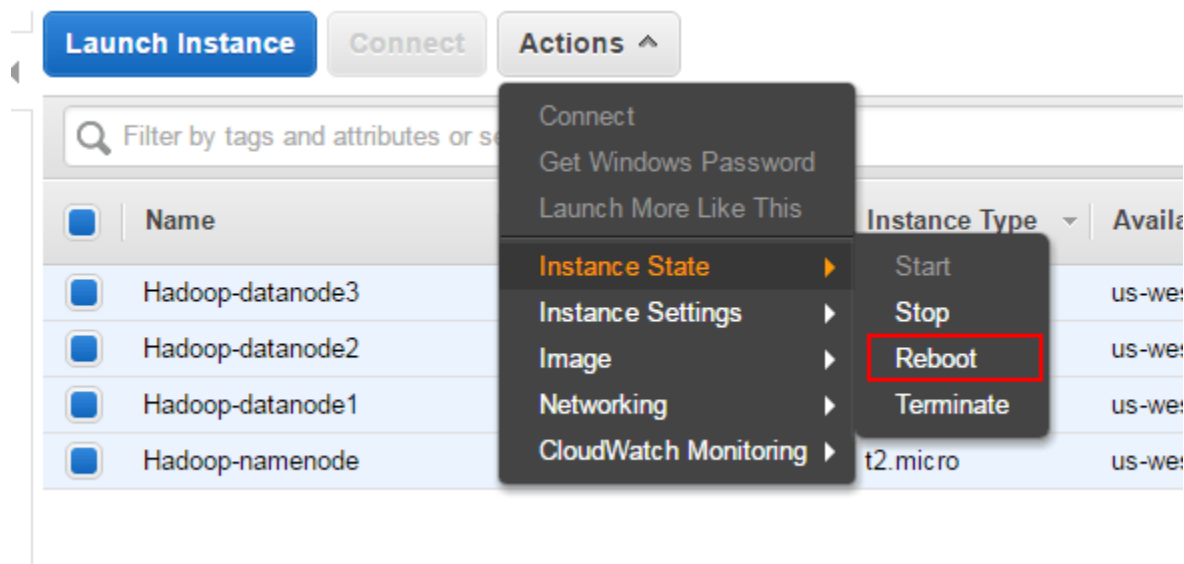
Some setup is common to all the nodes: NameNode and DataNodes. This is covered in this section.

### All Nodes: Update the instance

Let us update the OS with latest available software patches.

```
sudo apt-get update && sudo apt-get -y dist-upgrade
```

After the updates, the system might require a restart. Perform a Reboot from the EC2 Instances page.



## All Nodes: Install Java

Let us now install Java. We install the package: `openjdk-8-jdk-headless` on all the nodes.

```
sudo apt-get -y install openjdk-8-jdk-headless
```

## All Nodes: Install Apache Hadoop

Install Apache Hadoop 2.7.3 on all the instances. Obtain the link to download from the [Apache website](#) and run the following commands. We install Hadoop under a directory `server` in the home directory.

```
mkdir server

cd server

wget <Link to Hadoop 2.7.3>

tar xvzf hadoop-2.7.3.tar.gz
```

## All Nodes: Setup JAVA\_HOME

On each of the nodes, edit `~/server/hadoop-2.7.3/etc/hadoop/hadoop-env.sh`.

Replace this line:

```
export JAVA_HOME=${JAVA_HOME}
```

With the following line:

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

## All Nodes: Update `core_site.xml`

On each node, edit `~/server/hadoop-2.7.3/etc/hadoop/core-site.xml` and replace the following lines:

```
<configuration>

</configuration>
```

with these (as mentioned above, replace `<nnode>` with NameNode's public DNS):

```
<configuration>

  <property>

    <name>fs.defaultFS</name>

    <value><nnode>:9000</value>

  </property>

</configuration>
```

## All Nodes: Create Data Dir

HDFS needs the data directory to be present on each node: 1 name node and 3 data nodes. Create this directory as shown and change ownership to user ubuntu.

```
sudo mkdir -p /usr/local/hadoop/hdfs/data

sudo chown -R ubuntu:ubuntu /usr/local/hadoop/hdfs/data
```

## Configuring NameNode

After performing configuration common to all nodes, let us now setup the NameNode.

### Namenode: Password Less SSH

As mentioned before, we need password-less SSH between the name nodes and the data nodes. Let us create a public-private key pair for this purpose on the namenode.

```
namenode> ssh-keygen
```

Use the default (`/home/ubuntu/.ssh/id_rsa`) for the key location and hit enter for an empty passphrase.

## Datanodes: Setup Public Key

The public key is saved in `/home/ubuntu/.ssh/id_rsa.pub`. We need to copy this file from the namenode to each data node and append the contents to `/home/ubuntu/.ssh/authorized_keys` on each data node.

```
datanode1> cat id_rsa.pub >> ~/.ssh/authorized_keys  
  
datanode2> cat id_rsa.pub >> ~/.ssh/authorized_keys  
  
datanode3> cat id_rsa.pub >> ~/.ssh/authorized_keys
```

## Namenode: Setup SSH Config

SSH uses a configuration file located at `~/.ssh/config` for various parameters. Set it up as shown below. Again, substitute each node's Public DNS for the HostName parameter (for example, replace `<nnode>` with EC2 Public DNS for NameNode).

```
Host nnode
```

```
  HostName <nnode>
```

```
  User ubuntu
```

```
  IdentityFile ~/.ssh/id_rsa
```

```
Host dnode1
```

```
HostName <dnode1>
```

```
User ubuntu
```

```
IdentityFile ~/.ssh/id_rsa
```

```
Host dnode2
```

```
HostName <dnode2>
```

```
User ubuntu
```

```
IdentityFile ~/.ssh/id_rsa
```

```
Host dnode3
```

```
HostName <dnode3>
```

```
User ubuntu
```

```
IdentityFile ~/.ssh/id_rsa
```

At this point, verify that password-less operation works on each node as follows (the first time, you will get a warning that the host is unknown and whether you want to connect to it. Type *yes* and hit enter. This step is needed once only):

```
namenode> ssh nnode
```

```
namenode> ssh dnode1
```

```
namenode> ssh dnode2
```

```
namenode> ssh dnode3
```

## Namenode: Setup HDFS Properties

On the NameNode, edit the following

file: `~/server/hadoop-2.7.3/etc/hadoop/hdfs-site.xml`

Replace:

```
<configuration>

</configuration>
```

With:

```
<configuration>

  <property>

    <name>dfs.replication</name>

    <value>3</value>

  </property>

  <property>

    <name>dfs.namenode.name.dir</name>

    <value>file:///usr/local/hadoop/hdfs/data</value>

  </property>

</configuration>
```

## Namenode: Setup MapReduce Properties

On the NameNode, copy the file

(~/server/hadoop-2.7.3/etc/hadoop/mapred-site.xml.template) to  
(~/server/hadoop-2.7.3/etc/hadoop/mapred-site.xml). Replace:

```
<configuration>

</configuration>
```

With this (as above replace <nnode> with NameNode's public DNS):

```
<configuration>

  <property>

    <name>mapreduce.jobtracker.address</name>

    <value><nnode>:54311</value>

  </property>

  <property>

    <name>mapreduce.framework.name</name>

    <value>yarn</value>

  </property>

</configuration>
```

## **Namenode: Setup YARN Properties**

Next we need to set up ~/server/hadoop-2.7.3/etc/hadoop/yarn-site.xml on the NameNode. Replace the following:

```
<configuration>
```

```
<!-- Site specific YARN configuration properties -->

</configuration>
```

With (as before, replace `<nnode>` with NameNode's public DNS):

```
<configuration>

  <!-- Site specific YARN configuration properties -->

  <property>

    <name>yarn.nodemanager.aux-services</name>

    <value>mapreduce_shuffle</value>

  </property>

  <property>

    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>

    <value>org.apache.hadoop.mapred.ShuffleHandler</value>

  </property>

  <property>

    <name>yarn.resourcemanager.hostname</name>

    <value><nnode></value>

  </property>
```



```
</configuration>
```

## Namenode: Setup Master and Slaves

On the NameNode, create `~/server/hadoop-2.7.3/etc/hadoop/masters` with the following (replace `<nnode>` with the NameNode's public DNS):

```
<nnode>
```

Also replace all content in `~/server/hadoop-2.7.3/etc/hadoop/slaves` with (replace each of `<dnnode1>`, etc with the appropriate DataNode's public DNS):

```
<dnnode1>
```

```
<dnnode2>
```

```
<dnnode3>
```

## Configuring Data Nodes

After covering configuration common to both NameNode and DataNodes, we have a little bit of configuring specific to DataNodes. On each data node, edit the file `~/server/hadoop-2.7.3/etc/hadoop/hdfs-site.xml` and replace the following:

```
<configuration>
```

```
</configuration>
```

With:

```
<property>
```

```
<name>dfs.replication</name>
```

```
<value>3</value>

</property>

<property>

  <name>dfs.datanode.data.dir</name>

  <value>file:///usr/local/hadoop/hdfs/data</value>

</property>
```

## Starting the Hadoop Cluster

After all that configuration, it is now time to test drive the cluster. First, format the HDFS file system on the NameNode:

```
namenode> cd ~/server

namenode> ./hadoop-2.7.3/bin/hdfs namenode -format
```

Finally, startup the Hadoop Cluster. After this step you should have all the daemons running on the NameNode and the DataNodes.

```
namenode> ./hadoop-2.7.3/sbin/start-dfs.sh

namenode> ./hadoop-2.7.3/sbin/start-yarn.sh

namenode> ./hadoop-2.7.3/sbin/mr-jobhistory-daemon.sh start historyserver
```

Check the console carefully for any error messages. If everything looks OK, check for the daemons using `jps`.

```
namenode> jps
```

You can also check on the datanodes for java processes.

```
datanode1> jps
```

## Check the Web UI

Once the cluster is running, we can check the web UI for the status of the data nodes. Go to `<nnode>:50070` for the Web UI and verify that the 3 data nodes added are online.

## Summary

This article showed you how to setup Apache Hadoop on an Amazon EC2 cluster. Starting with standard Ubuntu 16.04 LTS instances, we configured the machines with Java and Apache Hadoop. Next we covered setting up various Hadoop components according to the role on each node. In the next article in this series on Apache Hadoop, we look into running jobs on the cluster.

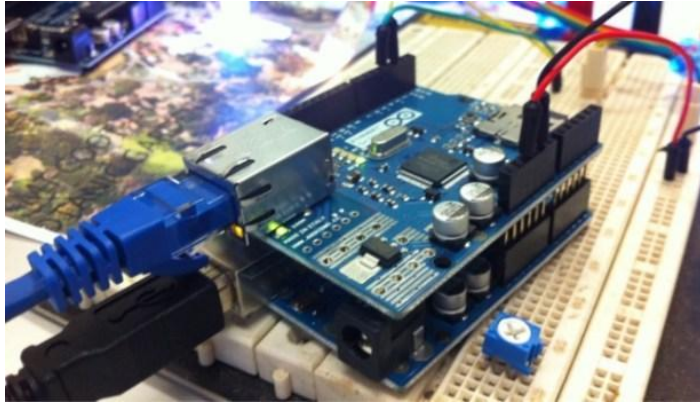
See Also



[Apache Spark - Setup Cluster on AWS](#)



[Hadoop Tutorial](#)



## [Setting up Apache Web Server with an SSL Certificate](#)

Author [Jay Sridhar](#) Posted on [March 30, 2017](#) Categories [Uncategorized](#) Tags [apache](#), [hadoop](#)

### Post navigation

PREVIOUS Previous post: [Java Regex – Simple Patterns](#)

NEXT Next post: [Apache Spark – Setup Cluster on AWS](#)

### POPULAR



[Using AES for Encryption and Decryption in Python Pycrypto](#)



[How to Generate RSA Keys in Java](#)



[Python How to Check if File can be Read or Written](#)



[Parsing XML in Python](#)

- 



[How to Setup an Apache Hadoop Cluster on AWS EC2](#)

- 



[Converting Between XML and JSON Using JAXB and Jackson](#)

- 



[Using AES With RSA for File Encryption and Decryption in Java](#)

- 



[Java - Reading a Large File Efficiently](#)

- 



[Using Jackson for JSON Serialization and Deserialization](#)

- 



[How to Read CSV File in Java](#)

## **LATEST**

- [Java 9 Modules Tutorial – Getting Started](#)
- [How to Generate Bitcoin Addresses in Java](#)
- [How to use Docker to Deploy Jupyter with Nginx](#)
- [Run Python Web Application in Docker using Nginx and uWsgi](#)
- [Nginx Inside Docker – Website Root Configuration](#)
- [Nginx Inside Docker with Ubuntu 16.04 HOWTO](#)
- [Python Regular Expressions Tutorial – Part 2](#)
- [Python Regular Expressions Tutorial – Part 1](#)
- [Python Crypto Basics for Building a Blockchain](#)
- [Using SSL to Secure Tomcat with HTTPS](#)

- 
- [Home](#)
- 

[Novixys Software Dev Blog](#) [Proudly powered by WordPress](#)

